

Digital Console — L2 Review Series

Three sessions on three consecutive days, starting today · plus a short stand-up on the days between

This is not training. L2 are the reviewers. By Friday we need three decisions:

1. Are the numbers **true**?
2. Are the alerts **right**?
3. Is the console **useful** in a real incident?

L1 workshops start only after this closes. If that means a delay, the delay is the right call — L1 will believe whatever the console tells them, so it has to be right before they see it.

Why L2 first

L1 can't spot a wrong number

They'll quote it to a customer or close a case on it. You can tell.

You know what normal looks like

Semati failing 55% on a quiet Tuesday. Activation codes that mean nothing vs everything. That judgement isn't in the code — it needs to get into the rules.

A week of real use beats any demo

Between sessions you work real cases in it. Everything wrong, noisy or missing gets logged and triaged the next morning.

Schedule

	Session	Focus	Ends with
Day 1	1 · Is the data true?	Lineage, freshness, every definition on the dashboard, attribution, coverage	12 definitions signed off
<i>Day 1→2</i>	<i>stand-up</i>	Triage what you logged	fixes started same day
Day 2	2 · Are the alerts right?	Rule mechanics, all 43 rules, baselines, INC0016809, correlation & escalation	every threshold has a verdict

Day 2→3	stand-up	Triage	
Day 3	3 · Is it useful?	Troubleshoot board, timeline, three real incidents replayed, ownership	verdict + ranked backlog
Day 4–5	live use	You run it on real traffic	Friday: ranked list of what's wrong

Fill in actual dates and times before sharing. Each session includes a break at the halfway point.

Session 1 — Is the data true?

#

1.1
1.2
1.3
1.4
1.5

—

1.3 is the centre of the whole week. Every box on the order-flow tree is a definition you can veto. Some already have open questions:

- Pending → **stuck** currently needs a gateway commit response present. Is 3h the right cut-off?
- **Not Assigned** now excludes partner-fulfilled orders (tygo/soob ship their own SIMs). Is the partner list complete?
- **Delivered** uses a fixed list of delivery states. Is that list current?
- **Activated** trusts `onboarding_orders.activated`. Does BSS agree?

Homework: take five real cases you handled that day. Find each in the console. Log every place it disagreed with what you knew, or couldn't answer the question at all.

Session 2 — Are the alerts right?

#

2.1
2.2
2.3
2.4

2.5

—

Current state: **43 rules · 33 metrics · 15 P1 · 24 P2 · 4 P3**, and five rules disabled (dealer_activity_drop, hyperpay_hard_down, ownership_fail_spike, payment_zatca_gap, tap_hard_down) — we decide together whether they stay off. Note tap_hard_down + Tap reconciliation both being off means a total Tap outage is currently invisible — first item on the session-2 agenda.

Every rule gets one of four verdicts:

Verdict

KEEP

CHANGE

DELETE

MISSING

The four new BSS rules need your judgement most. I set those thresholds from a single incident; you've seen hundreds.

Homework: for every alert that fires before we next meet, mark it **actionable** or **noise**. That one column is how we tune for the next six months.

Session 3 — Is it useful?

#

3.1

3.2

3.3

3.4

3.5

—

You drive. Simulate replay against the real incident window, no hindsight. The test for each: **would this have saved you time on the actual bridge call — and if not, what exactly was missing?**

A polite “yes it's useful” is worth nothing. The useful answer is the specific click that wasn't there, the number you had to go elsewhere for.

Between sessions — the loop

You, during the day: use it on a real case — something looks off — log it — keep working.

Us, next morning: triage → classify → fix or schedule → you see the fix.

Finding type	Example	Who fixes	Target
Data wrong	a KPI disagrees with BSS/prod	console owner	same day — worst kind
Alert noisy	fires nightly, never actionable	L2 + owner together	next session
Alert missing	an incident nothing caught	console owner	next session
Hard to use	six clicks to answer one question	console owner	backlog, ranked in S3

Log everything in L2-Review-Log.xlsx. Ask in the channel, not in a DM — the channel is the record.

What I need from you

Be specific, not polite. “The payment page is confusing” can’t be fixed. “Decline reason isn’t visible without two clicks” was fixed the same day.

Report numbers that look wrong first. Every serious defect so far was found by someone saying “that looks wrong” — the reseller figures, the staleness banner, the stuck payments. All of them.

Say when an alert wasted your time. Noise isn’t a minor complaint; it’s how a monitoring system dies. A rule you’ve learned to ignore is worse than one that doesn’t exist.

By Friday I want a ranked list of what’s wrong — not a list of what’s good.

After this week

Area	Owner	Cadence
Daily health — data fresh and sane?	L2 on duty	each morning
Alert tuning — thresholds, noise, gaps	L2 lead + console owner	weekly, charts open
New detections after an incident	whoever ran the incident	written up within 48h
Definitions — what a KPI counts	L2 agree, owner implements	on change
Training L1	L2, once this review closes	next series

Green light for L1 is a decision we make in session 3, not an assumption.

Before session 1

Facilitator:

```
ssh yosri@172.31.38.152 && sudo su -  
cd /apps/console/server && set -a; . ../.env; set +a  
node /apps/console/server/postdeploy-check.cjs # 22 checks — all green before  
you present
```

- Accounts created for every attendee at **L2 Digital** role (05-USER-SETUP-RUNBOOK.md)
- Pick today's live examples for 1.3 — a scenario dies if the example returns “not found”
- Teams links filled into workshop.config.json

Attendees: bring a real case from this week that was harder than it should have been.

Materials

File

Salam-Console-L2-Review.pptx

L2-REVIEW-PLAN.md

L2-Review-Log.xlsx

03-SCENARIOS.md

05-USER-SETUP-RUNBOOK.md

docs/INC0016809-detection-gap.md